

hw4

May 15, 2024

1 Support Vector Machine Algorithm

```
[ ]: # Importing Libraries
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt

from sklearn.svm import SVC
from sklearn.utils import resample
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import GridSearchCV
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score
```

1.1 Loading train and test datasets

Test dataset

```
[ ]: test_data = pd.read_csv('./data/SalaryData_Test.csv')
print(test_data.head(10))
```

	age	workclass	education	educationno	maritalstatus	\
0	25	Private	11th	7	Never-married	
1	38	Private	HS-grad	9	Married-civ-spouse	
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	
3	44	Private	Some-college	10	Married-civ-spouse	
4	34	Private	10th	6	Never-married	
5	63	Self-emp-not-inc	Prof-school	15	Married-civ-spouse	
6	24	Private	Some-college	10	Never-married	
7	55	Private	7th-8th	4	Married-civ-spouse	
8	65	Private	HS-grad	9	Married-civ-spouse	
9	36	Federal-gov	Bachelors	13	Married-civ-spouse	

	occupation	relationship	race	sex	capitalgain	\
0	Machine-op-inspct	Own-child	Black	Male	0	
1	Farming-fishing	Husband	White	Male	0	
2	Protective-serv	Husband	White	Male	0	

3	Machine-op-inspct	Husband	Black	Male	7688
4	Other-service	Not-in-family	White	Male	0
5	Prof-specialty	Husband	White	Male	3103
6	Other-service	Unmarried	White	Female	0
7	Craft-repair	Husband	White	Male	0
8	Machine-op-inspct	Husband	White	Male	6418
9	Adm-clerical	Husband	White	Male	0

	capitalloss	hoursperweek	native	Salary
0	0	40	United-States	<=50K
1	0	50	United-States	<=50K
2	0	40	United-States	>50K
3	0	40	United-States	>50K
4	0	30	United-States	<=50K
5	0	32	United-States	>50K
6	0	40	United-States	<=50K
7	0	10	United-States	<=50K
8	0	40	United-States	>50K
9	0	40	United-States	<=50K

Train dataset

```
[ ]: train_data = pd.read_csv('./data/SalaryData_Train.csv')
print(train_data.head(10))
```

	age	workclass	education	educationno	maritalstatus \
0	39	State-gov	Bachelors	13	Never-married
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse
2	38	Private	HS-grad	9	Divorced
3	53	Private	11th	7	Married-civ-spouse
4	28	Private	Bachelors	13	Married-civ-spouse
5	37	Private	Masters	14	Married-civ-spouse
6	49	Private	9th	5	Married-spouse-absent
7	52	Self-emp-not-inc	HS-grad	9	Married-civ-spouse
8	31	Private	Masters	14	Never-married
9	42	Private	Bachelors	13	Married-civ-spouse

	occupation	relationship	race	sex	capitalgain \
0	Adm-clerical	Not-in-family	White	Male	2174
1	Exec-managerial	Husband	White	Male	0
2	Handlers-cleaners	Not-in-family	White	Male	0
3	Handlers-cleaners	Husband	Black	Male	0
4	Prof-specialty	Wife	Black	Female	0
5	Exec-managerial	Wife	White	Female	0
6	Other-service	Not-in-family	Black	Female	0
7	Exec-managerial	Husband	White	Male	0
8	Prof-specialty	Not-in-family	White	Female	14084
9	Exec-managerial	Husband	White	Male	5178

	capitalloss	hoursperweek	native	Salary
0	0	40	United-States	<=50K
1	0	13	United-States	<=50K
2	0	40	United-States	<=50K
3	0	40	United-States	<=50K
4	0	40	Cuba	<=50K
5	0	40	United-States	<=50K
6	0	16	Jamaica	<=50K
7	0	45	United-States	>50K
8	0	50	United-States	>50K
9	0	40	United-States	>50K

1.2 Data understanding

1.2.1 Dataset size

```
[ ]: print(
    f"Test dataset size: {test_data.shape[0]}\n"
    f"Train dataset size: {train_data.shape[0]}"
)
```

Test dataset size: 15060
Train dataset size: 30161

```
[ ]: print(train_data.describe())
print(train_data.info())
```

	age	educationno	capitalgain	capitalloss	hoursperweek
count	30161.000000	30161.000000	30161.000000	30161.000000	30161.000000
mean	38.438115	10.121316	1092.044064	88.302311	40.931269
std	13.134830	2.550037	7406.466611	404.121321	11.980182
min	17.000000	1.000000	0.000000	0.000000	1.000000
25%	28.000000	9.000000	0.000000	0.000000	40.000000
50%	37.000000	10.000000	0.000000	0.000000	40.000000
75%	47.000000	13.000000	0.000000	0.000000	45.000000
max	90.000000	16.000000	99999.000000	4356.000000	99.000000

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 30161 entries, 0 to 30160

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	age	30161 non-null	int64
1	workclass	30161 non-null	object
2	education	30161 non-null	object
3	educationno	30161 non-null	int64
4	maritalstatus	30161 non-null	object
5	occupation	30161 non-null	object
6	relationship	30161 non-null	object
7	race	30161 non-null	object

```

8   sex          30161 non-null  object
9   capitalgain  30161 non-null  int64
10  capitalloss  30161 non-null  int64
11  hoursperweek 30161 non-null  int64
12  native       30161 non-null  object
13  Salary       30161 non-null  object
dtypes: int64(5), object(9)
memory usage: 3.2+ MB
None

```

1.2.2 Data Description

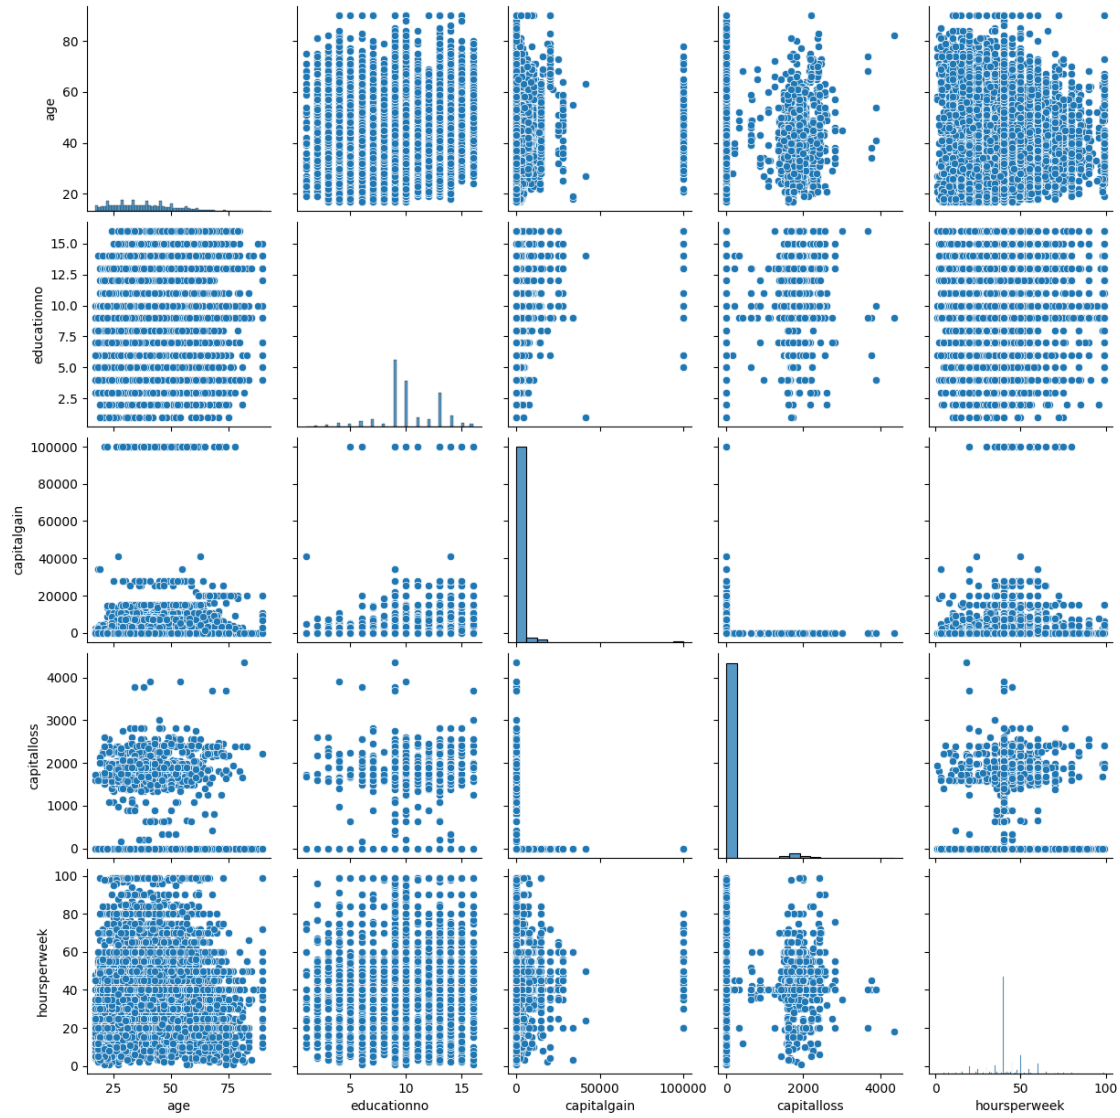
Term	Definition
Education	Highest level of education
workclass	Which business sector a person works for
educationno	Years spent in school?
maritalstatus	Relationship status
occupation	Current occupation
relationship	If they have a spouse or child
race	A social construct to group people
sex	Gender
capitalgain	Profit gained from the sale of property or investment
capitalloss	Value lost from the seller when selling property or investment at a lower cost
hourseperweek	Hours worked per week
native	A person's country of citizenship
Salary	Category on whether a person makes less than or greater than or equal to \$50,000

1.3 Data Visualization

1.3.1 Data Scatterplot

```
[ ]: sns.pairplot(train_data)
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x137ca67b0>
```




```
[ ]: sns.pairplot(train_data, hue='sex')
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x30960c920>
```



```
[ ]: sns.pairplot(train_data, hue='workclass', y_vars=['Salary'])
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x30c1d8770>
```



1.4 Data Processing

1.4.1 Label Encoder

In order for a model to work with the data provided, especially strings, the data must be categorized with discrete values.

Setting `train_data` columns with `object` values into enumerated integers.

```
[ ]: le = LabelEncoder()
train_data['Salary'] = le.fit_transform(train_data['Salary'])
train_data['workclass'] = le.fit_transform(train_data['workclass'])
train_data['education'] = le.fit_transform(train_data['education'])
train_data['maritalstatus'] = le.fit_transform(train_data['maritalstatus'])
train_data['occupation'] = le.fit_transform(train_data['occupation'])
train_data['relationship'] = le.fit_transform(train_data['relationship'])
train_data['race'] = le.fit_transform(train_data['race'])
train_data['sex'] = le.fit_transform(train_data['sex'])
train_data['native'] = le.fit_transform(train_data['native'])
```

1.5 Model Building

1.5.1 Preparing the train and test values

```
[ ]: x_train = train_data
y_train = train_data['Salary']
x_test = test_data
y_test = test_data['Salary']

# Perform label encoding
x_test['education'] = le.fit_transform(x_test['education'])
x_test['relationship'] = le.fit_transform(x_test['relationship'])
y_test = le.fit_transform(y_test)

print(f"x_train shape: {x_train.shape}\n"
      f"y_train shape: {y_train.shape}\n"
      f"x_test shape: {x_test.shape}\n"
      f"y_test shape: {y_test.shape}")
```

```
x_train shape: (30161, 14)
y_train shape: (30161,)
x_test shape: (15060, 14)
y_test shape: (15060,)
```

1.5.2 Factors that I believe affect salary the most

I've chosen the following factors that affect salary the most: - Education - A person's level of education gives the access to higher paying opportunities. - Relationship - Dependents may or may not affect a person's ability to work more

```
[ ]: from sklearn import svm
      from sklearn.inspection import DecisionBoundaryDisplay

      x_train = x_train[['education', 'relationship']].copy()
      x_test = x_test[['education', 'relationship']].copy()
```

1.5.3 Training the data

Training with polynomial kernel causes the whole program to freeze

```
[ ]: for kernel in ("linear", "rbf"):
      clf = svm.SVC(kernel=kernel, gamma=2)
      clf.fit(x_train, y_train)

      print(
          f"Accuracy of kernel {kernel}: {accuracy_score(y_test, clf.
↪predict(x_test)) * 100:.2f}%"
      )
```

Accuracy of kernel linear: 75.43%

Accuracy of kernel rbf: 81.59%

1.5.4 Plotting model boundary

```
[ ]: x_min, x_max = train_data['education'].min() - 1, train_data['education'].max()
      ↪+ 1
      y_min, y_max = train_data['relationship'].min() - 1, train_data['relationship'].
      ↪max() + 1

      fig, ax = plt.subplots(figsize=(6,4))
      ax.set(xlim=(x_min, x_max), ylim=(y_min, y_max))

      common_params = {
          "estimator": clf,
          "X": x_train,
          "ax": ax
      }

      DecisionBoundaryDisplay.from_estimator(**common_params,
                                          response_method="predict",
                                          plot_method='pcolormesh',
                                          alpha=0.3)
      DecisionBoundaryDisplay.from_estimator(**common_params,
                                          response_method="decision_function",
                                          plot_method="contour",
                                          )

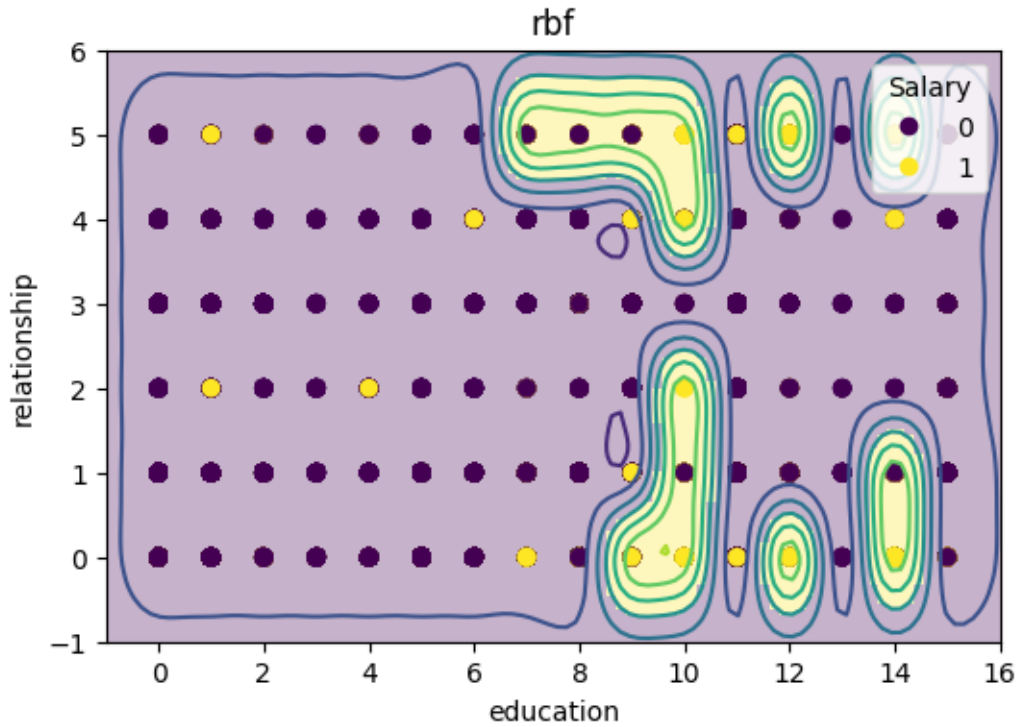
      scatter = ax.scatter(
          train_data['education'], train_data['relationship'], c=y_train
```



```

)
ax.legend(*scatter.legend_elements(), loc="upper right", title="Salary")
ax.set_title(kernel)
plt.show()

```



[]: